# DeltaSieve: Targeted Visual Validation of Spreadsheet Transformations

Alexander Aue-Johr
*Institute of Rural Studies*
*Thünen-Institute*
Braunschweig, Germany
*Harz University of Applied Sciences*
Doctoral Candidate, Germany
ORCID 0009-0001-8683-3630

***Abstract*—Native visual validation of spreadsheet transformations is expensive when every worksheet must be rendered. DeltaSieve selects potentially affected regions using differences in cell values, effective style properties, layout, and dependencies of conditional formatting and charts. Unknown dependencies retain explicit capture fallbacks. Paired native rebuilds and retained transformation controls distinguish comparison evidence from changes introduced by Excel itself. The system records selected coverage, exact pixel observations, process measurements, and reusable evidence. In a first-stage comparison of five transformations across 27 workbooks from Data.gov, Destatis, and Enron, 102 of 135 planned comparisons produced reports with explicit coverage limits. Independent native controls illustrate both a removed heading fill and an inherited border extending beyond populated cells. These results characterize observed effects and coverage within the tested scope. They do not estimate corpus-wide detection recall.**




Original, SheetShrink, and XLStylesTool

Check Performance, Inquire, and NXPowerLite

Fig. 1. Sheet Atlas for the nine highest-ranked Data.gov workbooks in the first-stage snapshot. Red tiles mark recorded visual differences. Neutral tiles include regions without complete coverage. The enlarged Adage Fixed Factor reference shows a shortened title band in three transformations. The two image groups apply only to the matched reference range.

## I. Introduction

A spreadsheet transformation can preserve cell values and still change the visible document. Removing an empty formatted cell can expose a background or alter automatic row height. Changing a referenced input can affect a distant conditional-formatting region or chart. Checking only edited cell payloads therefore leaves important transformation effects unobserved.

Native rendering supplies an operational way to observe these effects in the spreadsheet application that users employ. Exhaustive rendering, however, can be expensive for large sheets and corpora. Stored worksheet extents may also include empty or uniformly formatted tails. Dividing these extents by a fixed cell grid can create many images despite very little finite visible content. The challenge is to select useful rendering work while retaining explicit uncertainty about unsupported features and incomplete comparisons.

Figure 2 summarizes the selection and evidence workflow. DeltaSieve combines package-level change analysis with native visual comparison. It follows effective style identities across table renumbering, considers layout and dependency effects, and records the reasons for its capture choices. Supported conditional-formatting and chart dependencies can establish that an unchanged input cannot have changed the associated presentation under the supported model. Changed or unresolved dependencies remain capture candidates. The native stage measures geometry, renders bounded image batches, and compares retained evidence under a recorded state-pair policy.

The scientific question concerns the detection and cost of this selection. A dependency-aware planner is useful only if its reduced work retains the relevant effects in the evaluated domain. We therefore separate selector coverage, pixel observation and defect attribution. Independent full-reference samples and mutations test detection. Repeated same-input observations test rendering variability. A changed image remains an observation whose cause may require further evidence.

Existing spreadsheet dependency slicing motivates careful treatment of indirect effects. Existing workbook comparison tools also provide reports of structural and cell-level changes. DeltaSieve's evaluation must establish the additional contribu-

tion of its native-rendering selection and evidence workflow. An interactive Observatory and portable report make those observations inspectable. They support the method's evaluation without replacing its independent detection tests.

DeltaSieve grew out of the need to investigate visual regressions in SheetShrink, a workbook cleanup system described in a companion manuscript. That manuscript studies cleanup rules and preservation structures. Here we study the selection, capture and evidence needed to evaluate rendered changes across workbook transformations, including alternative cleanup tools.

## II. Preservation and Comparison Contract

### A. Inputs, states, and observations

DeltaSieve validates a pair of spreadsheet states produced by a transformation. The package representation supplies evidence about where visible changes could occur. Native Excel rendering supplies the image observations for selected regions. The two sources of evidence serve different roles. A candidate is a reason to inspect a region; pixel inequality is an observation that still needs attribution.

The fast comparison uses Original and transformed workbooks after a full rebuild and save in their existing, separately owned Excel processes. This gives both partners the native recalculation step. Stored and ordinary-saved XML controls remain available for follow-up. Their changes explain potential causes without automatically enlarging the rebuilt-pair capture plan. An independently requested Check Performance comparison has its own transformed state and coverage.

### B. Relevant visual area and candidate coverage

The relevant area includes values, formulas, effective visible formatting, text overflow, and objects. A union across relevant input states is needed to cover content removed by a transformation. Empty trailing geometry has a separate contract. The capture plan must retain enough evidence to distinguish a change inside the relevant area from a larger or smaller exterior canvas.

Candidate selection considers changed cell properties and layout, as well as regions affected through conditional formatting and graphical dependencies. Style and shared-string references require comparison of their effective contents because saving can renumber the corresponding tables. Supported reference reasoning can exclude a dependent region when its rule and inputs are unchanged. A changed dependency makes its affected region a candidate. Unresolved reference semantics require explicit fallback coverage.

A selected-region comparison does not by itself demonstrate that every relevant change in the workbook was selected. We evaluate that property using independently rendered reference areas and controlled mutations. The evaluation must record which reference cases completed and which exceeded their budgets. The paper's coverage claims are bounded by these observations and the supported dependency model.

### C. Verdicts, failures, and reusable evidence

A successful comparison requires usable evidence from both partners. Missing hashes, failed native commands, and timeouts remain unsuccessful outcomes. Heartbeats establish add-in reachability and do not establish calculation or capture progress. Durable operation receipts and phase measurements support investigation of stalls.

The raw pixel and geometry observations remain attached to a finding after review. Product regressions, rendering variation, environment-dependent values, and verified capture-extent differences have separate classifications. A volatile formula path explains a possible cell-value change while other changed pixels in the same image can require further analysis. No global pixel tolerance is assumed.

Reusable observations bind the input states, relevant configuration, producing validation implementation, and declared rendering environment. Package-content identity is evaluated independently of ZIP entry order and compression. Reusing a completed study preserves its source manifest and compatibility checks. Portable reports retain the observations and their provenance for inspection without rerunning Excel.

## III. Method

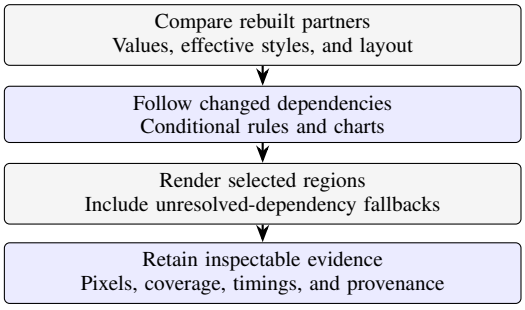



Fig. 2. DeltaSieve separates dependency-aware selection, native observation, and evidence. A proven unchanged dependency can exclude rendering. An unresolved dependency retains its recorded fallback. Targeted coverage remains explicit.

### A. Comparison subject and observation model

Figure 3 separates package analysis, native observation and outcome accounting. The input is an original/transformed pair, a state-preparation policy and a renderer environment. The principal comparison opens separate owned Excel processes and observes both partners after full rebuilding in place. Ordinary-save and pre-rebuild controls retain the effects of serialization and calculation. Conversion of a legacy XLS source precedes this subject and has its own failure outcome.

Let $A$ and $B$ denote the compared states, $E$ their justified relevant visual area and $C$ the selected observation regions. The intended selector obligation is to include every transformation-induced visual difference inside $E$ in $C$. Pixel equality on $C$ tests the selected observations. It does not establish that the selector met that obligation. Independent reference captures and visible mutations evaluate this second

**Require:** State pair $(A, B)$, observation policy $P$
**Ensure:** Scoped comparison evidence and explicit outcomes
Record input identities and observation policy
Resolve value, effective-style and layout changes
Follow supported conditional and chart dependencies
Include conservative scopes for unresolved effects
Open owned native workers and prepare states under $P$
**for all** bounded packages of selected regions **do**
    Measure native geometry in both states
    Record zero-area outcomes and partition visible areas
    Render paired tiles within pixel and memory budgets
    Compare complete compatible image receipts
    Retain differences, equality evidence and failures
**end for**
Record selected coverage and unfinished work
Release owned resources and record confirmed exits

Fig. 3. DeltaSieve's logical comparison stages. Missing or failed captures retain explicit outcomes. Independent reference controls evaluate selection.

question. Renderer variability and unsupported dependencies remain explicit boundaries of the model.

### *B. State preparation and attribution controls*

Opening a workbook can repair or normalize its stored representation, saving can remove structure, and recalculation can change cached results. The comparison therefore names the operations that produced each observed state. Let $W_0$ be the raw package, $W_s$ its ordinary-save control and $W_r$ the rebuilt saved control. A native optimization comparison additionally retains $W_{cp}$ after Check Performance and Save As, and $W_{cpr}$ after rebuilding that result. The transformed workbook has its own corresponding controls.

The primary observation compares independently prepared rebuilt partners. Control differences explain which changes arose while opening, saving, optimizing or rebuilding. Ordinary Save and Save As remain distinct generation recipes. One cannot substitute a cheaper recipe under an existing control name and retain its evidence meaning. Legacy format conversion and explicit repair also produce separately identified artifacts.

Each control records its parent identity, native environment and generation recipe. A reusable control artifact saves preparation work but supplies no new visual verdict. A comparison still needs compatible observations of both partners. If opening or a required operation fails, that state remains unavailable. The workflow does not manufacture its image from another state. This separation also permits attributing a reduced package to Excel's save behavior without crediting that reduction to the tested cleanup operation.

### *C. Streaming change evidence*

The package analyzer uses borrowed XML byte views with bounded hierarchy and attribute state. It resolves shared strings, cell formats and relationships before interpreting a coordinate change. String and formula interpretation can require decoding, but worksheet analysis does not instantiate a general XML DOM or an Excel object per cell. This is important on bloated inputs with many empty formatting records.

**Keep cell evidence outside the heap**

Abbreviated Rust serialization

```
let mut record = [0u8; RECORD_BYTES];
record[..4].copy_from_slice(
    &cell.row.to_le_bytes());
record[4..6].copy_from_slice(
    &cell.column.to_le_bytes());
// Two flag bytes occupy positions 6-7.
record[8..].copy_from_slice(&cell.digest);
output.write_all(&record)?;
```

A 40-byte disk record stores row, column, two flag bytes and a 32-byte digest. Ordered records support a streaming comparison without retaining one heap object per cell.

This is the serialized layout, not the Rust struct layout. Style tables, changed cells and native images still need separate memory.

Fig. 4. The fixed-size cell evidence stream. Coordinates and flags accompany the content/style digest. Abbreviated Rust.

Cell evidence is written as fixed-size records containing coordinates, a SHA-256 fingerprint and classification flags. The current record is 40 bytes. The record stream resides in an owned scratch file instead of a heap map of cell objects. Sorted streams are merged by coordinate, distinguishing changed, added and removed cells. Equal stream digests avoid repeated cell comparison. The merge is linear in the combined record counts, excluding XML parsing, style resolution, hashing and selected dependency work. Scratch I/O is an explicit cost of this memory organization. Shared strings, styles, layout records, selected changes and native images remain additional memory consumers.

Figure 4 shows the serialized evidence layout. The writer requires unique ascending numeric coordinates, rejects nested or unclosed cells, and hashes a cell's text incrementally. A very large cell body therefore need not become one temporary string. Record readers require complete records. A truncated scratch file cannot become a shorter valid sheet. This strict stream contract permits a two-way merge without a separate sorting pass for supported ordered worksheets.

The fingerprint preserves distinctions that the model has not justified merging. Unsupported serialization can therefore produce conservative extra work rather than a semantic-equivalence claim. Canonicalizing style-table references is a specific supported operation, not a general normalization of arbitrary formulas, XML extensions or renderer behavior. The same byte-processing substrate serves content comparison, supported reference scans and rebuild-eligibility facts, while their decision predicates remain separate.

A format index has no cross-package meaning. Style identity resolves fonts, fills, borders, number formats and base-format inheritance, so a permutation of format tables does not create a change at every cell. Changed shared context,

**Follow the changed input**

Abbreviated dependency illustration

```
<!-- Only B1 changes: 0 to 1. -->
<conditionalFormatting sqref="D2:D40">
 <cfRule type="expression" dxfId="0">
  <formula>$B$1&gt;0</formula>
 </cfRule>
</conditionalFormatting>
<!-- The rule is unchanged. -->
<!-- D2:D40 still needs observation. -->
```

Changing B1 can change the rendered fill in D2:D40 even when all cells and rule XML in that area are unchanged. The application area becomes a native capture candidate.

This illustration assumes a supported rule and unchanged differential style. Unknown formulas retain fallback coverage.

Fig. 5. A changed precedent can require observation outside the edited coordinate. Abbreviated XML.

including themes or unsupported style representations, retains the conservative path. A removed blank cell can expose an inherited row or column format even when values are unchanged. Conversely, an absent untyped blank and an empty untyped record can supply the same blank input to a conditional rule. The full visual fingerprint and the rule-input fingerprint therefore have different responsibilities.

### *D. Dependency-aware candidate selection*

Direct value, formula, style and layout changes create candidates. Effects propagate beyond their edited coordinates. A row-height change can move later content, a column-width change can move content to its right, and a merged or centered span can extend the required text context. Changed inherited row formats are coalesced into contiguous runs, with unchanged rows splitting the runs. This reduces candidate objects while preserving the union of requested regions.

Conditional-format selection resolves the supported rule formula, application area and referenced inputs. A changed precedent selects the affected application region even when its own appearance is unchanged. Local rules, area-wide inputs and explicitly referenced inputs require different dependency scopes. Exclusion requires unchanged supported rules, their relevant input fingerprints and layout. Unknown names, dynamic references, extensions or unresolved scope retain fallback coverage. DeltaSieve does not evaluate the rule to predict Excel's final color. The native renderer supplies that observation.

Figure 5 shows an unchanged rule whose precedent changes. Selection follows the dependency to its application region. It does not infer that unchanged serialized cells inside that region imply unchanged pixels.

### *E. Invert supported rule references*

Enumerating every target cell of a whole-column conditional rule would defeat the purpose of selective rendering. For supported A1 references, DeltaSieve instead maps each changed source coordinate back to dependent target intervals. A relative scalar reference at row $a$ in a rule anchored at row $h$ reads source row $a + t - h$ when applied at target row $t$. A changed source row $x$ therefore selects target row

$$t = x - a + h.$$

An absolute reference to row $a$ selects all permitted target rows when $x = a$, and none otherwise. Columns follow the corresponding rule. Range endpoints can independently mix absolute and relative axes. The implementation solves the resulting axis intervals, combines row and column intervals and intersects them with the application and comparison extents. It includes both shared and per-area anchor interpretations conservatively for disjoint application areas.

This inversion avoids materializing the entire worksheet grid. Its work still depends on changed inputs, supported reference expressions and application areas. Cross-sheet references require a unique supported source sheet. Dynamic, external or ambiguous references retain fallback. Exhaustive small-grid regression controls compare the inverse result with direct forward reference evaluation, including mixed axes and reversed range endpoints. These tests verify the supported mapping, not arbitrary Excel formula semantics.

Charts have a separate dependency closure. Equal chart XML and cached series values are insufficient because source cells can change independently of the cache. The exclusion gate requires unchanged chart parts, supported source references, source values and styles, source layout and visibility, conditional dependencies, shared package context and complete equal native placement. Chart sources are inspected beyond the worksheet capture envelope. A source outside the visible table cannot disappear from this dependency check. Missing partners, ambiguous names, user-shape dependencies and unsupported references retain chart capture. A conservative mode exports every chart.

Graphical objects require both package identity and native geometry. Relationship mapping binds an object to its actual worksheet. Supported lines include their stroke envelope. Zero geometric height alone cannot exclude visible ink. Cached previews of embedded objects can render even when Excel labels their shape kind unsupported. Unsupported objects therefore prevent a finite-bound claim unless their own explicit qualification succeeds.

### *F. Finite observation envelopes*

The relevant envelope includes values, formulas, supported visible formatting, text context and qualified object extents from both compared states. Removed content therefore remains in scope. Native-qualified drawing bounds refine the rebuilt-state envelope. Save-control changes retain their own attribution records without automatically enlarging that distinct rebuilt comparison.

TABLE I
SELECTION OBLIGATIONS. AN EXCLUSION APPLIES ONLY WITHIN THE SUPPORTED REPRESENTATION AND THE RECORDED STATE PAIR.

| Effect | Required context |
|---|---|
| Cell formatting | Resolved cell, row and column components |
| Row/column layout | Native effective geometry and later content |
| Conditional rule | Rule, application region and supported input closure |
| Chart | Source cells, package graph, formatting and live placement |
| Drawing or OLE | Resolved owner, frame, ink or preview and native bounds |
| Uniform tail | Complete transitions and a finite relevant envelope |

An empty exterior row cannot move earlier content. Effective-height selection and image selection use the same lower union boundary, while earlier empty rows remain relevant. Native zero-width and zero-height ranges keep their geometry and a non-rendered outcome. A transition to a visible range remains a finding. No white image is invented for a missing or zero-area observation.

Repeated inherited backgrounds need a finite-bound argument. For supported borderless solid fills, the analyzer retains explicit cells, row and column transitions, contextual neighbors and supported conditional application areas. Uniform unfilled column styles beneath a solid row background introduce no additional horizontal fill edge. Fonts and widths alone do not create such an edge in an empty band. Visible column fills, borders, unknown styles and unsupported rules prevent that exclusion. Supported finite filters and conditional regions remain in the envelope with following context. Full-grid or unresolved application areas retain the conservative path.

These conditions permit less rendering without sampling the selected union. They are relative to the supported model. An unresolved outlying object can still force a full-grid plan and exhaust the available rendering budget.

### G. All-state adaptive geometry

Candidate coalescing preserves holes and worksheet identity. Independent ranges are measured in batches through the native adapter. Every live state contributes to image-size qualification. A range hidden in one state can be large in its partner, so any oversized state forces further splitting. Each child is checked again. Large hidden spans can yield zero-area leaves rather than thousands of identical image requests. Figure 6 shows the batch completeness and geometry checks.

**Every live state controls the split**

Abbreviated Rust implementation

```
let measured = measure(&batch)?;
ensure!(
    measured.len() == batch.len(),
    "Incomplete geometry batch coverage"
);
for (candidate, geometries) in
    batch.into_iter().zip(measured)
{
    for geometry in geometries.values() {
        geometry.validate()?;
    }
    // Split if any live state is too large.
}
```

The planner measures the same candidate in every live state. A hidden range in one workbook can be visible in its partner. Any oversized state forces further partitioning.

A missing measurement is an error. Zero-area leaves retain an observation. A single oversized cell fails instead of disappearing.

Fig. 6. Every live state participates in geometry qualification. Abbreviated Rust.

The planner respects both area and edge budgets. A conservative point-to-pixel allowance guides partitioning, but returned image dimensions are independently validated. A single cell exceeding the budget fails instead of being omitted. Measured row and column geometry handles hidden or irregular spans that make a fixed cell-count tile unreliable. Later states can repartition an area and backfill earlier observations.

Internal boundaries retain required neighboring context for overflow, merged text and other supported layout effects. The final interior is cropped without resampling. A context region can change rendering even when its own pixels are later cropped. The renderer controls therefore compare a whole region with contextual parts and check seam consistency. Agreement of one seam is a local control, not a monolithic oracle for the workbook.

### H. Native image qualification

Office.js requests images through `Range.getImage()`. Batching amortizes cross-runtime synchronization. Image results are consumed only after the pending synchronization completes, as shown in Figure 7. The adapter bounds both queued image count and aggregate pixels. Geometry, encoding, transient buffers and upload windows have separate limits.

A successful API response does not prove complete capture. The worker verifies expected dimensions, budget compliance and crop geometry before accepting the image. Returned PNG bytes can first be hashed to avoid uploading identical payloads. Unequal encodings still require decoded-pixel comparison, because PNG serialization is not visual identity. Exact canonical RGBA pixels supply the comparison, with no registration, resampling or general color tolerance. Image receipts bind command, region, state and slot identity, preventing reordered or incomplete upload batches from becoming comparisons of unrelated images.

Raw geometry and origins remain attached to the observation. An exterior-size classification can admit an unpaired right or bottom extension only when each pixel exactly repeats the common boundary. A zero-area counterpart requires an entirely opaque-white retained extent for the corresponding

**Queue, synchronize, then consume**

Abbreviated JavaScript implementation

```
const images = chunk.map(tile =>
  worksheet.getRange(
    CaptureGeometry.address(tile.raw)
  ).getImage()
));
await context.sync();
return images.map(image => image.value);
```

One synchronization completes the queued Office.js image requests. Returned images are then dimension-checked, cropped without resampling and finalized before another bounded batch.

The batch pixel limit also applies. A heartbeat while synchronization is pending does not prove that a capture completed.

Fig. 7. Office.js results are consumed after synchronization. Abbreviated JavaScript.

empty-extent classification. These derived classifications retain the original observations. They do not erase changed overlapping content or establish a global tolerance.

*I. Attribution after a finding*

A native difference can originate in the transformation, calculation, capture context or the environment. DeltaSieve keeps the observation while diagnostics examine those hypotheses. The formula tracer follows supported A1 references, absolute axes, quoted sheet names, ranges and translated shared formulas from changed saved values. It searches for paths to random/time or environment-dependent functions. Cycles, unsupported expressions and bounded-search exhaustion remain explicit. The search budget limits explanation work, not capture coverage.

A path to a volatile function establishes a dependency, not exclusive causation. Other changed pixels in the same region can have another cause. Source cached values also cannot stand in for values saved from the compared rebuilt states. The report distinguishes their provenance. Repeated original/original controls and unchanged-package controls test renderer behavior independently of the transformation. Counterfactual restoration of a suspected removed witness tests a narrower causal mechanism. None of these operations increases pixel tolerance or automatically clears an unrelated finding.

*J. Execution, omission and evidence reuse*

Each comparison group owns its Excel processes, private desktop and WebView profile. Office calls are serial within a workbook. Independent groups can progress concurrently under memory admission. Cleanup distinguishes requested termination from confirmed process exit. Setup, opening, geometry, capture, comparison and persistence have separate timing boundaries.

A finite-work watchdog advances after verified work, including validated geometry responses and completed image comparisons. Heartbeats establish reachability while a native call is pending, but do not extend a deadline by themselves. Durable commands and receipts reconcile reconnects. Repeated delivery cannot restart an already completed operation or silently duplicate a pending one. Failed attempts retain produced evidence and receive no equality verdict.

Some observations can be omitted before native work. A direct byte comparison can establish that the original and transformed inputs are identical. ZIP metadata alone does not authorize that omission. A separate rebuild-eligibility check reuses dependency facts from the package scan and a strict part/feature allowlist. Formulas, refresh sources, unknown extensions, duplicate members and unresolved metadata keep rebuilding enabled. Cached values or an empty capture plan do not establish calculation independence.

Completed-result reuse requires compatible source identities, selected parts, implementation, state-preparation policy and rendering environment. Static report-viewer assets are excluded from the renderer recipe. Uncompressed member digests make identity independent of ZIP order and compression. A prior part selection must cover every newly requested component. User-selected weaker component policies remain explicit evaluation choices, with correspondingly bounded claims. Interrupted runs and missing image receipts cannot supply completed evidence.

The cache admission condition can be expressed as compatibility of four independent facts. The previous observation must be complete, its required package contents must match, its selected component set must cover the new request, and its renderer/preparation policy must remain compatible. Matching only a filename or producer binary is insufficient. Matching workbook bytes without the renderer contract is likewise insufficient for a native result. A changed compression level can preserve member identity, while a new ordinary save can change it and invalidate admission.

Generation caches, completed-comparison caches and partial working evidence have different roles. An interrupted attempt may retain useful images for inspection, but its missing coverage cannot be supplied by relabeling the attempt complete. Publication records retain the original result and later classification rather than overwriting the evidence after a retry. This is important in a corpus campaign where several machines or attempts contribute observations.

Portable reports preserve input identities, observation regions, exclusions, raw findings, attribution and resource measurements. The Observatory supports inspection and replay without becoming a source of visual truth. Detailed transport schemas, UI controls, experiment launchers and historical failed attempts belong to the versioned artifact. They do not change the selection or pixel-observation contract described here.

### K. Machine-Accessible Campaign Monitoring

The source inventories described in Section IV motivate monitoring across long-running campaigns. They contain 34,065 content-addressed Enron table candidates, 9,117 Data.gov XLSX files and 1,182 Destatis XLSX files. Candidate inventory size, successful preparation and completed comparison coverage remain separate quantities.

The control room exposes a local HTTP/JSON interface for compact run summaries, worker and workbook state, comparison reports and retained evidence. Bounded summaries preserve active attempts while limiting completed-item detail. Run and attempt identities prevent observations from earlier executions from being reported as current work. Process liveness, heartbeats and the time of last verified progress are represented separately, allowing an external monitor to distinguish a reachable process from a progressing attempt. A WebSocket stream provides bounded snapshots and incremental status updates for live clients.

These interfaces supported local agent workflows during overnight runs. Agents could inspect stalled attempts and record suspected visual differences for subsequent author review. Capture receipts, pixel comparisons and retained images supply the validation evidence. Agent annotations do not establish visual equality or convert incomplete coverage into a completed result. Explicit control requests and command-status queries support intervention without requiring an agent to infer state from the rendered dashboard.

## IV. Evaluation

RQ1 asks whether the selector captures the known visible effects in the tested mutation classes and independent reference sample. RQ2 measures how much native capture work and elapsed time it avoids. RQ3 examines the reproducibility and investigability of retained findings, including failures and rendering variation.

### A. Corpus identity and outcome accounting

The evaluation treats Data.gov, Destatis and Enron as distinct populations. Archive checksums, member names and duplicate occurrences identify each selected corpus version. Legacy XLS conversion is reported separately from XLSX cleanup. A historical Enron development selection contains 20,872 inputs, with 20,503 prepared artifacts and 369 unsuccessful preparation outcomes at its recorded preparation checkpoint. Its selected visual scope contains 16,799 workbooks. These development counts do not establish completed visual coverage. The replacement EDRM Enron v2 XML handover contains 34,065 distinct candidate SHA-256 identities totaling 8.470 decimal GB of uncompressed payload. Manifest checks verified transport while retaining original bytes and upstream extraction provenance. The frame includes legacy Excel and other OLE candidates. Native format recognition, XLSX conversion and visual eligibility are separate outcomes, so the candidate count does not measure workbook coverage.

Microsoft Graph OLE objects require an additional eligibility distinction. Optional native PowerPoint conversion can recover their chart data into newly generated XLSX workbooks. Those workbooks have a different preparation history from successfully opened original Excel files and are reported as a separate derived-data population. Successful recovery or equality of cell values does not establish visual equivalence of the source Graph chart or coverage of its parent presentation. They do not contribute visual-equality observations without a separately executed native comparison.

TABLE II
Frozen XLSX source inventories. Sizes are decimal GB of listed workbook files and do not measure completed native coverage.

| Corpus | XLSX members | XLSX total (GB) |
|---|---|---|
| Data.gov | 9,117 | 25.273 |
| Destatis | 1,182 | 2.364 |

The two frozen XLSX inventories in Table II come from the versioned source profiles. Each listed archive member has an individual SHA-256 identity. Their source-frame sizes are separate from the eligible, completed and independently referenced subsets.

A dated evaluation cutoff will report compared workbooks, differences, failures, unsupported representations and unfinished work. Repeated attempts have their own counts. An unchanged package can bypass visual work under an exact input-identity check, but a failed capture cannot contribute an equality observation. The full recursive Enron XML-derived corpus is a separately identified population.

### B. Sampling and reference controls

Seeded samples use frozen source frames and size strata chosen before visual inspection. Development stress cases are labelled separately. Independent reference plans bypass candidate selection. Timeouts and unsupported inputs remain in the selected denominator. Controlled edits require an observed native effect or an explicit invariance outcome.

A reconciled historical Enron development sample retained twelve seeded source occurrences. Seven completed native references over 33 declared rectangles, two failed image-size preflight and three required unsupported graphical or embedded-object envelopes. All 33 completed rectangles had identical overlapping RGBA pixels. Two had additional opaque-white exterior strips, totaling 196,657 pixels, which were retained as raw extent differences. This sample is not held out, and failed or unsupported references remain in its denominator.

A separate sample selected one input per size quartile from 30,301 prepared Enron XLSX files using a fixed seed before inspecting selected outputs. All four native references qualified over 41 stored-dimension rectangles, with matching image dimensions and zero changed RGBA pixels. Prior development exposure remains unclassified. The ordinary v15 selector rendered 14 images totaling 78.0 million pixels, compared with 82 reference images totaling 398.0 million pixels. All seven targeted image pairs were equal. One byte-identical input pair required no targeted images. These are render-work counts

under distinct plans, with adaptive targeted partitions and fixed 256-row reference tiles.

A fixed-seed Data.gov sample selected one member per size quartile from the 9,117-entry inventory before joining cleanup outcomes. Two initial references failed size preflight. Smaller independent partitions subsequently completed those inputs, retaining the failed attempts. All four references cover 188 areas with equal dimensions and pixels. Targeted selection used eleven equal pairs. Prior exposure is unclassified; this all-equal sample cannot estimate recall on changed inputs.

A separate twelve-occurrence Destatis sample used the same fixed-seed, size-quartile selection policy. All twelve selected packages contain graphical or embedded parts and remain unsupported by the current dimension-only reference preparer. They were not replaced by easier inputs. An optional native-object qualification extension admitted one of the same twelve inputs over two fixed areas. Its image rectangle fit inside the cover-sheet area, and both reference areas had identical overlapping RGBA pixels. The data-sheet image had a 20-pixel-high opaque-white exterior strip totaling 25,820 pixels. The remaining eleven inputs retained unsupported outcomes for legacy drawings, embedded objects or drawing effects. This extension checks native object rectangles, with explicit rejection of unsupported types and out-of-area objects. It does not certify decorative ink beyond those rectangles or establish corpus recall.

### C. Detection and cost metrics

A completed independent reference determines whether a mutation has a visible effect under the recorded observation policy. A selector miss occurs when such an effect falls outside selected coverage. Detection is measured per mutated workbook pair and per mutation class. A partially covered effect retains a separate outcome even when another difference makes the pair detectable. Pixel coverage requires a documented mapping between reference and selected capture coordinates, including extent differences.

Selection overhead counts candidates whose completed paired images are equal. This measures avoidable work within the tested policy. Pixel findings retain their attribution as transformation effects, rendering variation or unresolved differences. Failed references, missing partner images and timeouts retain explicit outcomes. Report both completed-reference detection and the fraction of all selected cases with usable reference evidence.

### D. Mutation strata and statistical scope

Table III summarizes nine designed controls with exact reference-mask coverage. INDIRECT used the unsupported-rule fallback. Dependency resolution and broader unresolved-dependency strata remain open. The defined-name case changes only the range binding, rebuilding its unchanged SUM formula from 3 to 6. The color-scale and data-bar cases change rule colors.

Both states are rebuilt and saved. NOW, RAND, INDIRECT and OFFSET require reference-policy controls. Raw pixels and calculation context support attribution, alongside identical-input rendering controls.

TABLE III
DESIGNED MUTATIONS AFTER NATIVE REBUILD. PIXELS COUNT CHANGED RGBA LOCATIONS IN INDEPENDENTLY FIXED REFERENCES. EVERY SELECTED MASK EXACTLY MATCHED ITS REFERENCE MASK AT THEIR COMMON A1 ORIGIN. THESE CONTROLS DO NOT ESTIMATE POPULATION RECALL.

| Mutation | Pairs | Reference pixels | Selected area |
|---|---|---|---|
| Wrapping | 1 | 1,775 | A1:D3 |
| Merge A1:D1 | 1 | 3,873 | A1:D3 |
| INDIRECT conditional fill | 2 | 5,070 each | A1:AA3 |
| Formula errors | 2 | 366 / 358 | A1:D3 |
| Color scale | 1 | 5,727 | A1:D3 |
| Data bar | 1 | 4,304 | A1:D3 |
| Defined-name binding | 1 | 53 | A1:D3 |

For zero misses in $n$ independent trials with common miss probability, the exact one-sided 95% upper bound is $1 - 0.05^{1/n}$. Workbook mutations share a cluster. Population estimates require strata and inclusion probabilities, with development cases reported separately.

### E. Matched cost accounting

Cost accounting separates phases, images, pixels, retained bytes, elapsed time and coverage unions. Overlapping captures count as repeated work. Matched trials control preparation, environment, image budgets and reuse through open, rebuild and save. Repeated pairs measure variation. Timeouts remain censored. Corpus projections require geometry-calibrated strata and hidden-region/object-extent assumptions. Production queues cannot isolate tuning effects.

### F. Opening, batching and independent Excel pairs

A bounded opening experiment used four copies of a synthetic workbook with eight million-row, three-column sheets on one Windows host. Two reversed-order rounds measured median batch times of 118.388 seconds sequentially and 41.714 seconds in independent processes sharing a private desktop. The endpoint includes Office.js readiness. Staging and teardown are excluded, and cache and background-load effects remain. This supports overlapping startup for that configuration, not arbitrary concurrent COM calls.

A warmed-up 64-range geometry experiment returned identical measurements in scalar and batched modes. The scalar passes each took 16.094 seconds and the batch passes 0.250 seconds. Twelve-region image packages, including four zero-area regions, took 3.125 seconds at requested size one and 0.579 seconds at size eight, with identical decoded image dimensions and pixels. Opening was excluded and 250-ms observation polling quantizes these timings. The experiment establishes a transport-amortization opportunity under its recorded workload rather than a corpus throughput factor.

Earlier independent-capture trials repeatedly lost one concurrent add-in to its 120-second startup deadline. A later paired experiment used one host and shared inactive desktop with independently owned pair states. Over three repetitions per completed configuration, moving from one pair to two increased

capture-only throughput from 3.752 to 7.030 pairs/minute with binary PNG transport and from 7.590 to 12.312 with hash-first transport. Peak observed job memory rose from 1,525 to 2,685 MiB and from 1,472 to 2,666 MiB, respectively. Every tile in this selection was equal, favoring the hash-first path. Requested configurations for three through eight pairs were rejected before launch by RAM admission. Those rejections are not zero-throughput measurements or an Excel scaling ceiling.

The historical image-budget protocol selected the smallest budget within 5% of the fastest eligible sum of per-scope median workloads. Both 18-case trials selected 32 MP. The replication sums were 245.14, 235.70 and 230.59 seconds for 16, 32 and 64 MP. Thus 64 MP was fastest and had no observed budget-block failures. Historical local-host, VBA/GDI+ and VSTO/Excel-DNA probes answer separate transport, rasterization or geometry questions. They do not rank interchangeable screenshot engines.

### *G. Coverage qualification before speed*

A synthetic million-row request returned an Office.js PNG without an API error in 2.58 seconds. It contained only $85 \times 32{,}766$ pixels against the required 20,971,521-pixel height. The boundary probe found 1,638 rows complete and 1,639 incomplete on that installation. This demonstrates why completion and dimensional qualification are separate conditions. An original-only one-column Lysimeter sweep subsequently covered its requested 978,624 rows in 576 images and 814.6 seconds after readiness. It supplied hashes and dimensions, not an original/transformed visual verdict.

The historical 62-case protocol completed 48 cases and retained 14 failures. Four unresolved drawing-bound worksheet outcomes account for repeated failures. They are not 14 independent crashes. Six of eight coverage workbooks completed, covering 462 of 466 requested worksheets. The failed observations remained incomplete instead of receiving a smaller undocumented range or an equality classification.

### *H. Reducing work while retaining the selected union*

The FRPP workbook has 307,920 rows, including 304,140 hidden rows. Of 19,248 original grid requests, 18,872 have zero area. A read-only geometry probe drove the Rust adaptive planner to 64 leaves after 710 probes in 2.815 seconds. This is a planner observation, not a completed paired Office.js comparison. A separate native visibility-transition control exposed a row only in the transformed state. It retained the original's zero-area observation and reported rendering and geometry differences without claiming paired-pixel comparison there.

A 208-worksheet diagnostic exposed 5,192,711 intermediate candidates. Grouping contiguous changed row formats reduced that count to 603,267 with identical content envelopes. Both diagnostics stopped before complete rasterization. This establishes reduced intermediate representation size and no measured end-to-end speedup. It motivates representing shared layout effects as intervals.

Uniform inherited backgrounds exposed a related native-work amplification. One Destatis pair had 128,912 versus 2,197 column definitions. The former plan measured 16,101 regions, although most were outside the explicit-cell envelopes. The guarded finite-background plan produced 165 image pairs, all identical in a fresh comparison. Selection and transport took 41.371 seconds and comparison 1.379 seconds. The earlier attempt was interrupted after more than 40 minutes. A second Enron pair reduced a former 65,664-region plan to 73 measured regions, all with identical images. Its selection and transport took 28.114 seconds. The latter case retained an empty finite AutoFilter and all serialized-cell extent. These are selected development observations with censored earlier attempts, not matched speedup estimates. Regressions retain distant fill transitions, borders, overrides and unsupported rule contexts.

Among 134 retained SheetShrink comparisons against the selected originals, median selection and transport took 7.4 seconds on Destatis (42 pairs), 30.4 on Data.gov (46), and 3.0 on Enron (46). Historical capture configurations vary. These timings describe recorded work, without estimating a speedup against full capture. On the Hospital Data workbook, 904 and 905 pre-grid candidates on one sheet expanded to four NXPowerLite and 3,297 Inquire regions respectively. Both comparisons were equal within their distinct scopes.

An inherited border can extend past every populated cell. In the Data.gov QT/RR workbook, the rebuilt dimensions of eleven affected sheets end at column L, but Inquire removes an inherited bottom border on row 1. An explicit reference at `XFB1:XFD3` found a 193-pixel line missing at the far edge, with equal measured geometry. Of 90,112 rectangles in the ordinary capture plan, 2,101 intersect columns A:L or row 1. The rebuilt Inquire files also declare a default column width of 9.140625 and set the custom-height flag, while retaining the original default row height of 12.75 points. This count describes a potential reduction in capture work. The two regions alone do not qualify a smaller plan because other layout and inherited appearance changes may remain.

Conversely, unsupported objects can dominate a conservative plan. Two retained attempts contained explanatory sheets responsible for 104,833 and 185,856 regions. After approximately 5.3 hours, neither full capture had completed. Removing two zero-height lines in a separately supplied reference area changed 1,042 RGBA pixels because their strokes remained visible. The long-running fallback therefore cannot be removed merely because an object's geometric height is zero.

### *I. Comparator pilot with independent native controls*

Five frozen synthetic pairs were supplied to xlCompare 12.5.0.0 and DiffEngineX 3.18.0.0 through their report commands. Both exports included an offscreen B2000 edit. Neither listed a hidden B3000 edit under the tested commands. Direct visual and sparse-formatting changes produced no exported cell differences. A dependent pair exported the changed B1 precedent without a separate conditional-fill or chart-rendering

TABLE IV
NATIVE CHART CONTROLS. IMAGE COUNT INCLUDES BOTH STATES. AN EXCLUDED CHART SUPPLIES NO NEW PAIRED-PIXEL OBSERVATION.

| Control | Images | Exclusions | Differences |
|---|---|---|---|
| Unchanged, deterministic | 0 | 1 | 0 |
| Unchanged, conservative | 2 | 0 | 0 |
| Changed source, deterministic | 2 | 0 | 1 |

finding. Edition and option coverage remains incomplete, including xlCompare's reported first-ten-rows export limit.

DeltaSieve completed the same five pairs using their ten hash-verified inputs. The identical pair was excluded by byte identity. Four pairs produced native pixel findings. The offscreen edit changed 47 overlapping pixels. Direct, dependent and sparse-formatting pairs changed 7,057, 3,950 and 1,281 grid pixels, respectively. The dependent pair also changed 36,505 chart pixels. The region containing hidden B3000 rendered equally. These observations distinguish edited cell inputs from their dependent visual effects.

An independently supplied reference plan bypassed candidate selection and content clipping. It used Data!A1:N30 for the short pairs, and continuous 256-row tiles through rows 3072 and 2304 for offscreen and sparse pairs. All five comparisons completed. The identical control remained equal. Only the B2000 and J2099 tiles changed in their respective long pairs, while the direct and dependent pairs also differed in their broader reference areas. Pair-level outcomes agree with the selected run. This is a bounded control, not an estimate of corpus recall or pixel-union coverage. Different chart and capture extents prevent pooling their raw pixel counts. No matched cross-tool speedup follows from these trials.

Two representation controls permuted cell-XF and shared-string indices while remapping their cell references. Independent `Data!A1:N30` renders were pixel-identical within both pairs. The style permutation selected zero regions. The shared-string permutation retained one equal image pair at `A1:D4`, exposing unnecessary capture work under the tested policy v15.

The wrapping and merge controls in Table III use independent `A1:N30` references. Merged A1:D1 includes populated cells and exercises native merge behavior.

Two additional comparator pairs contained 107 and 215 worksheets. The observed xlCompare times were 53.748 and 21.711 seconds. DiffEngineX required 131.692 and 161.221 seconds and represented every worksheet in its formula-mode reports. These single observations establish execution feasibility and report scope. No matched DeltaSieve cost trial was completed for those pairs.

### J. Chart dependency controls

A workbook created in Excel supplies three controlled chart cases. The unchanged pair differs only in document-description metadata. Deterministic selection excludes its one chart, while conservative mode renders both partners. A third pair changes a source value and retains chart capture. Table IV records the observed outcomes.

These controls test both sides of the exclusion gate. Equal caches alone are not the sufficient condition. The package and source-cell dependencies and native geometry must also agree. A separate rebuilt Lysimeter inventory contains 69 charts. Its gate excludes 34 and retains 35. Eighteen fallbacks have ambiguous duplicate names and 17 reference additional user-shape dependencies. This inventory did not complete a replacement image comparison and supplies no end-to-end runtime reduction.

### K. Attribution controls and retained counterexamples

Independent controls distinguish a transformation defect from capture or renderer behavior. In a fixed-height statistics sheet, the original-versus-original reference was pixel-identical, while the transformed pair repeatedly changed 25,640 pixels. Restoring 87 removed font-bearing blanks recovered the original image hash. A corrected SheetShrink output also matched over the frozen A1:L916 reference area. The case demonstrates a real witness omission rather than an excuse to loosen pixel comparison.

A migration control instead exposed inconsistent exterior context. Its relevant content ended at row 84, but capture included an empty row 85. That row changed from 12.75 points to hidden. A zero-width region then produced a height finding without an image. Applying the same union-content boundary to geometry and capture removed this exterior-context finding. A fresh run completed 38 regions, including 37 pixel comparisons, with no capture-region difference there. Separate height observations remain recorded.

Five Enron controls differed only in their ZIP comments before opening. Conservative capture compared ten charts. Eight were pixel-identical. Two charts, previously flagged in the SheetShrink comparisons, showed two and seven changed pixels in these controls with equal geometry. Rebuilt worksheet, style and chart parts were unchanged between control lanes. The controls demonstrate that native rendering can vary in these charts without a workbook-member transformation. Attribution of the original pixel masks remains unresolved, and the report retains the original findings alongside the separate controls.

Embedded-object mutations test another boundary. Moving a PDF-icon frame changed 6,473 RGBA pixels in a fixed reference area. Changing only its cached preview changed 5,242. Separately enlarging three zero-size text shapes produced 1,531, 1,737 and 1,487 changed pixels on three sheets, while removal controls were pixel-identical in their supplied areas. These cases show why package anchors, previews and visible ink need distinct treatment. They do not establish coverage for arbitrary OLE formats.

### L. What the measurements establish

Package bytes, candidate counts, native rectangles, rendered pixels and compared pixels are different units. A reduction in intermediate candidates can lower planner memory without changing the native image count. Adaptive partitioning can reduce images over hidden rows while preserving the same

TABLE V
FIRST-STAGE COVERAGE SNAPSHOT ON 11 NOVEMBER 2026. EACH CORPUS HAS 45 SLOTS. RAW AND NORMALIZED COUNT COMPLETED COMPARISONS WITH LIMITED COVERAGE. UNAVAILABLE COUNTS MISSING VARIANTS. NO ELIGIBLE COUNTS SLOTS WITHOUT A QUALIFYING WHOLE-WORKBOOK RESULT.

| Corpus | Raw | Norm. | Unavail. | Failed | No eligible | Active |
|---|---|---|---|---|---|---|
| Destatis | 27 | 2 | 1 | 14 | 1 | 0 |
| Data.gov | 38 | 0 | 0 | 4 | 1 | 2 |
| Enron | 37 | 1 | 4 | 3 | 0 | 0 |

selected cell union. Hash-first transport can reduce transmitted PNG bytes without avoiding Excel's image generation. The experiment ledger records these stages separately.

Similarly, process memory is not a single interchangeable counter. Native Excel, WebView2, the local host and the analysis process can hold different buffers. A sampled private-byte maximum can miss a short allocation peak, and a working-set maximum is not private commit. Scope and sampling interval accompany the reported memory. Phase durations can overlap and are not added as if they were independent. A claim about total campaign throughput includes startup, preparation and failed attempt costs under a declared reuse policy.

The independent reference plan remains distinct from those optimizations. It renders externally supplied areas without candidate selection, conditional-rule exclusion or production content clipping. Its author is responsible for the asserted coverage of those areas. A failed or oversized reference remains a failed reference case. Agreement between a selected plan and a reference over five deliberately constructed pairs supports those controls. It does not turn the same five pairs into a representative sample of three corpora.

### *M. Recorded coverage of the first comparison stage*

The first stage uses the nine workbooks with the largest absolute SheetShrink file-size reductions per corpus in the frozen selection. Each workbook contributes one slot per tool. Table V reports a timestamped snapshot of all 135 slots. The versioned snapshot identifies each rank and tool. The separate Enron ledger retains the detailed outcomes and baseline identities.

Of the three Enron failures, two followed Excel repair prompts and one required unresolved whole-grid coverage. The repaired Check Performance preparation and the completed SheetShrink comparison at rank two use different baseline hashes. The repaired preparation therefore provides no matched visual observation for that tool slot. These counts describe coverage of the selected high-reduction workbooks. Detection recall requires the independent reference experiments.

### *N. A border difference inside the common capture*

A developer-selected Enron control rendered `Tables!A1:I157` after inspecting a Check Performance finding. Four independent original captures were pixel-identical. SheetShrink and XLStylesTool preserved that reference exactly. Check Performance differed at 126,404 common pixels, while Inquire differed at 490 near the lower edge. Extending the reference through row 160 retained 127,256 and 916 changed common pixels, respectively. Inspection of the extended Inquire capture showed a missing lower horizontal table border inside the shared rectangle. Its rebuilt worksheet lacked original row 157 and its styled blank cells. The added context therefore separates this finding from a capture clipped at the table edge. NXPowerLite had no available output for this input. This post-hoc illustration covers one declared area and does not measure recall or establish whole-workbook preservation. The matching XLStylesTool result is retained alongside SheetShrink's result.

*a) Different capture coverage for the same original:* For one Data.gov PFAS workbook, the original package-member inventory matched across three tool comparisons. Inquire yielded 129 differing regions among 4,277 selected regions. An inspected range contained purple original pixels that were white in the candidate inside the common image extent. Check Performance yielded no differences in its 63 regions. XLStylesTool selected zero regions. These observations have different denominators and cannot rank the tools' visual preservation. The recorded selection and transport phases were 1,668.137, 24.025 and 0.032 seconds, respectively, and describe the validator's work on those selections.

Figure 1 shows the same `readme!A1:M22` reference area across all five tools. The original images agree across observations. Check Performance, Inquire and NXPowerLite omit the fill of D2:L2, changing the same 16,182 pixels. The two strips use the same crop coordinates. The reference manifest records source and image hashes. This localized result supplies no recall estimate.

### *O. Geometry without a visible surface*

Nine zero-width `LCCA` ranges differed in measured height (2000.25 versus 1827 points) without images. Version 16 retains their geometry but excludes these changes from visual findings after validating both zero surfaces. Missing images and visibility transitions remain independently checked. An unchanged-input native repeat completed 29 worksheet and 27 chart observations in 216 seconds with zero visual differences and nine nonvisual geometry observations. Coverage remains limited to the selected areas.

### *P. Operational outcomes and evaluation limits*

A Check Performance preparation campaign attempted 1,182 Destatis inputs and retained 1,016 outputs and 166 failures. Recovery diagnostics identified Name Conflict dialogs for 165 former opening timeouts, while one automation-property failure recurred. A separately retained manual success on another workbook shows that an automation failure does not establish tool incapability. These preparation outcomes remain separate from completed visual observations.

Shared-profile cleanup reproduced a peer taskpane restart, while isolated-profile controls passed. A watchdog corrected to count verified geometry responses completed 741 equal image pairs. Earlier failed receipts remain retained.

Bounded controls, stress cases and partial campaigns do not establish population recall, corpus-wide equivalence or matched end-to-end speedup. Broader independent references, mutation coverage and common-endpoint cost trials remain open. Manifests retain identities, revisions, settings and repeated attempts, with distinct complete, different, unsupported, failed and unfinished denominators.

## V. Discussion and Limitations

The central risk is an incomplete candidate model. A proof for supported inputs must state its dependencies and applicable namespace or rule subset. Unknown features need a capture fallback or an explicitly incomplete outcome. Runtime limits and failures reduce achieved coverage and must remain in the report.

A native renderer is also an experimental dependency. Recalculation, volatile formulas, file-location functions and rendering variability can change output without a transformation defect. A volatile dependency trace explains a supported path rather than every image difference. A white exterior canvas difference requires geometry and containment checks before classification. No global pixel tolerance is inferred from one rare anti-aliasing case.

The Observatory reduces the effort of investigating these outcomes. Portable reports and aggregate progress replay support inspection. Detailed process, sheet, capture and heartbeat replay remains open. History cannot establish authenticity through detail or supply unrecorded telemetry. Scientific claims rely on source identities, versions and independent evaluation.

## VI. Conclusion

DeltaSieve combines dependency-aware region selection with native rendering and inspectable evidence. The bounded experiments demonstrate concrete findings and fallbacks. Corpus-wide recall and matched capture cost remain open.


## Acknowledgment of AI Assistance

OpenAI ChatGPT and Codex assisted with drafting and editing the Abstract and Sections I–VI, literature discovery and organization, implementation and refactoring, test and experiment scripts, figure-generation code, and analysis of retained experiment evidence. This assistance included the method diagrams, XML examples and execution visualizations. Local agent workflows using these systems also supported overnight validation campaigns by inspecting status and logs, investigating stalled attempts, and documenting candidate visual regressions for author review.

The author directed development, reviewed source code, and performed manual testing and debugging. Agent observations supported investigation. Reported validation outcomes are grounded in the retained native captures and comparison records. The author remains responsible for the scientific claims, references, evidence selection, originality and final submitted content.